

Building Python Applications with Docker

Moshe Zadka

PyTexas 2017

Introduction

- ▶ Clarify terms

Introduction

- ▶ Clarify terms
- ▶ Common mistakes

Introduction

- ▶ Clarify terms
- ▶ Common mistakes
- ▶ Best practices

What are containers?

Share-less processes

What is Docker?

► client

What is Docker?

- ▶ client
- ▶ server (containerd)

What is Docker?

- ▶ client
- ▶ server (containerd)
- ▶ runner (runc)

What is Docker?

- ▶ client
- ▶ server (containerd)
- ▶ runner (runc)
- ▶ hub

What is a Container image?

- ▶ Tarball...

What is a Container image?

- ▶ Tarball...
- ▶ that looks like a linux root

What is a Container image?

- ▶ Tarball...
- ▶ that looks like a linux root
- ▶ Will also specify some metadata (e.g., entrypoint)

What is a Dockerfile?

Builds a Docker image

```
FROM ....
```

```
COPY from--context
```

```
RUN some command
```

```
...
```

Python application

catx/app.py

```
def hello_world(request):  
    return response.Response('California <3 Texas')  
  
with config.Configurator() as cfg:  
    cfg.add_route('hello', '/')  
    cfg.add_view(hello_world, route_name='hello')  
    app = cfg.make_wsgi_app()
```

catx/__main__.py

```
simple_server.make_server('0.0.0.0', 6543, app.app)
```

What's bad?

bad.docker

```
FROM debian:latest
RUN apt-get update
RUN apt-get -y install python3 python3-pip
RUN pip3 install pyramid
COPY catx /app/catx
ENV PYTHONPATH /app/
CMD python3 -m catx
```

What's wrong?

- ▶ latest

What's wrong?

- ▶ latest
- ▶ Layer explosion

What's wrong?

- ▶ latest
- ▶ Layer explosion
- ▶ Random system packages

What's wrong?

- ▶ latest
- ▶ Layer explosion
- ▶ Random system packages
- ▶ Random Python packages

What's wrong?

- ▶ latest
- ▶ Layer explosion
- ▶ Random system packages
- ▶ Random Python packages
- ▶ PYTHONPATH

What's wrong?

- ▶ latest
- ▶ Layer explosion
- ▶ Random system packages
- ▶ Random Python packages
- ▶ PYTHONPATH
- ▶ Shipping build environment

What's wrong?

- ▶ latest
- ▶ Layer explosion
- ▶ Random system packages
- ▶ Random Python packages
- ▶ PYTHONPATH
- ▶ Shipping build environment
- ▶ Using reference WSGI server

Base images

Choose foundation wisely (and stably)

Fork tagging

Renaming images for posterity

Fork tagging

```
IMAGE = 'pypy:3-slim '  
TAG = (datetime.datetime.now()  
        .isoformat().replace(':', '-')  
        .replace('.', '-'))  
  
NAME = 'moshez/pypy3:' + TAG  
cmd('docker', 'pull', IMAGE)  
cmd('docker', 'tag', IMAGE, NAME)  
cmd('docker', 'push', NAME)
```

Python Artifacts

sdist

- ▶ What `python setup.py sdist` creates
- ▶ Basically just a tarball of sources

Python Artifacts

wheel

- ▶ Fully built and ready
- ▶ Reliable installation

Python Artifacts

pex

- ▶ Python executable
- ▶ Installation is "copy"
- ▶ Does not work well with PyPy

Python Artifacts

Virtual environments

- ▶ Location specific
- ▶ Can be used with `dh_virtualenv`

What is Multistage Build?

- ▶ Build images sequentially

What is Multistage Build?

- ▶ Build images sequentially
- ▶ Throw away all except last

What is Multistage Build?

- ▶ Build images sequentially
- ▶ Throw away all except last
- ▶ Copy previous image

Multistage Build

```
FROM source as name1
...
FROM other-source
...
COPY —from=name1 ...
...
ENTRYPOINT [...]
```

Better Version 1

```
FROM moshez/pypy3:2017-10-30T09-29-03-882199 \  
as builder
```

Better Version 1

```
FROM moshez/pypy3:2017-10-30T09-29-03-882199 \  
as builder
```

```
RUN pypy3 -m venv /appenv
```

```
RUN /appenv/bin/pip install twisted pyramid
```

Better Version 1

```
FROM moshez/pypy3:2017-10-30T09-29-03-882199 \  
as builder
```

```
RUN pypy3 -m venv /appenv
```

```
RUN /appenv/bin/pip install twisted pyramid
```

```
COPY catx/ /mnt/src/catx/
```

```
COPY setup.py /mnt/src/
```

```
RUN /appenv/bin/pip install /mnt/src/
```

Better Version 1

```
FROM moshez/pypy3:2017-10-30T09-29-03-882199 \
as builder
```

```
RUN pypy3 -m venv /appenv
```

```
RUN /appenv/bin/pip install twisted pyramid
```

```
COPY catx/ /mnt/src/catx/
```

```
COPY setup.py /mnt/src/
```

```
RUN /appenv/bin/pip install /mnt/src/
```

```
RUN tar cvzf /appenv.tar.gz /appenv
```

Better Version 1

```
FROM moshez/pypy3:2017-10-30T09-29-03-882199 \  
as builder
```

```
RUN pypy3 -m venv /appenv
```

```
RUN /appenv/bin/pip install twisted pyramid
```

```
COPY catx/ /mnt/src/catx/
```

```
COPY setup.py /mnt/src/
```

```
RUN /appenv/bin/pip install /mnt/src/
```

```
RUN tar cvzf /appenv.tar.gz /appenv
```

```
FROM moshez/pypy3:2017-10-30T09-29-03-882199
```

```
COPY --from=builder /appenv.tar.gz /
```

```
RUN tar xvzf /appenv.tar.gz
```

```
ENTRYPOINT ["/appenv/bin/python", "-m", "twisted", \  
            "--wsgi=catx.app.app"]
```

Locking

```
$ git checkout -b updating-third-party
$ docker build -t temp-image -f harden.docker .
$ docker run --rm -it temp-image > requirements.str
$ git commit -a -m 'Update 3rd party'
$ git push
$ # Follow code workflow
```

Locking

```
FROM moshez/pypy3:2017-10-30T09-29-03-882199  
COPY harden-requirements requirements.loose.txt /  
ENTRYPOINT ["/harden-requirements"]
```


Locking

```
import subprocess, sys
subprocess.check_output([sys.executable, "-m", "venv",
                        "/envs/loose"])
subprocess.check_output(["/envs/loose/bin/pip", "install",
                        "/requirements/loose.txt"])
frozen = subprocess.check_output(["/envs/loose/bin/pip",
                                   "freeze", "--all"])
with open("/requirements/frozen.txt", 'wb') as fp:
    fp.write(frozen)
```

Summary

- ▶ Use multi-stage builds

Summary

- ▶ Use multi-stage builds
- ▶ Lock Python requirements

Summary

- ▶ Use multi-stage builds
- ▶ Lock Python requirements
- ▶ Think about your entrypoint